

```
In [94]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import Ridge, RidgeCV, Lasso
```

```
In [95]: df=pd.read_csv(r"C:\Users\pucha\Downloads\Advertising.csv")
df
```

Out[95]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
...	...	...	...	...
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

200 rows × 4 columns

```
In [96]: df.head()
```

```
Out[96]:
```

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

```
In [97]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype  
---  -
0   TV           200 non-null   float64
1   Radio        200 non-null   float64
2   Newspaper    200 non-null   float64
3   Sales        200 non-null   float64
dtypes: float64(4)
memory usage: 6.4 KB
```

```
In [98]: df.tail(12)
```

```
Out[98]:
```

	TV	Radio	Newspaper	Sales
188	286.0	13.9	3.7	20.9
189	18.7	12.1	23.4	6.7
190	39.5	41.1	5.8	10.8
191	75.5	10.8	6.0	11.9
192	17.2	4.1	31.6	5.9
193	166.8	42.0	3.6	19.6
194	149.7	35.6	6.0	17.3
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

```
In [99]: df=df[['TV','Radio','Sales','Newspaper']]  
df.columns=['tv','radio','sales','newspaper']
```

```
In [100]: df.head(10)
```

```
Out[100]:
```

	tv	radio	sales	newspaper
0	230.1	37.8	22.1	69.2
1	44.5	39.3	10.4	45.1
2	17.2	45.9	12.0	69.3
3	151.5	41.3	16.5	58.5
4	180.8	10.8	17.9	58.4
5	8.7	48.9	7.2	75.0
6	57.5	32.8	11.8	23.5
7	120.2	19.6	13.2	11.6
8	8.6	2.1	4.8	1.0
9	199.8	2.6	15.6	21.2


```
In [101]: df.describe()
```

```
Out[101]:
```

	tv	radio	sales	newspaper
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	15.130500	30.554000
std	85.854236	14.846809	5.283892	21.778621
min	0.700000	0.000000	1.600000	0.300000
25%	74.375000	9.975000	11.000000	12.750000
50%	149.750000	22.900000	16.000000	25.750000
75%	218.825000	36.525000	19.050000	45.100000
max	296.400000	49.600000	27.000000	114.000000

```
In [102]: sns.lmplot(x="radio",y="sales",data=df,order=2,ci=None)
```

```
Out[102]: <seaborn.axisgrid.FacetGrid at 0x20e1f8c9b40>
```



```
In [103]: df.fillna(method='ffill',inplace=True)
```

```
In [104]: x=np.array(df['radio']).reshape(-1,1)
          y=np.array(df['sales']).reshape(-1,1)
          df.dropna(inplace=True)
```

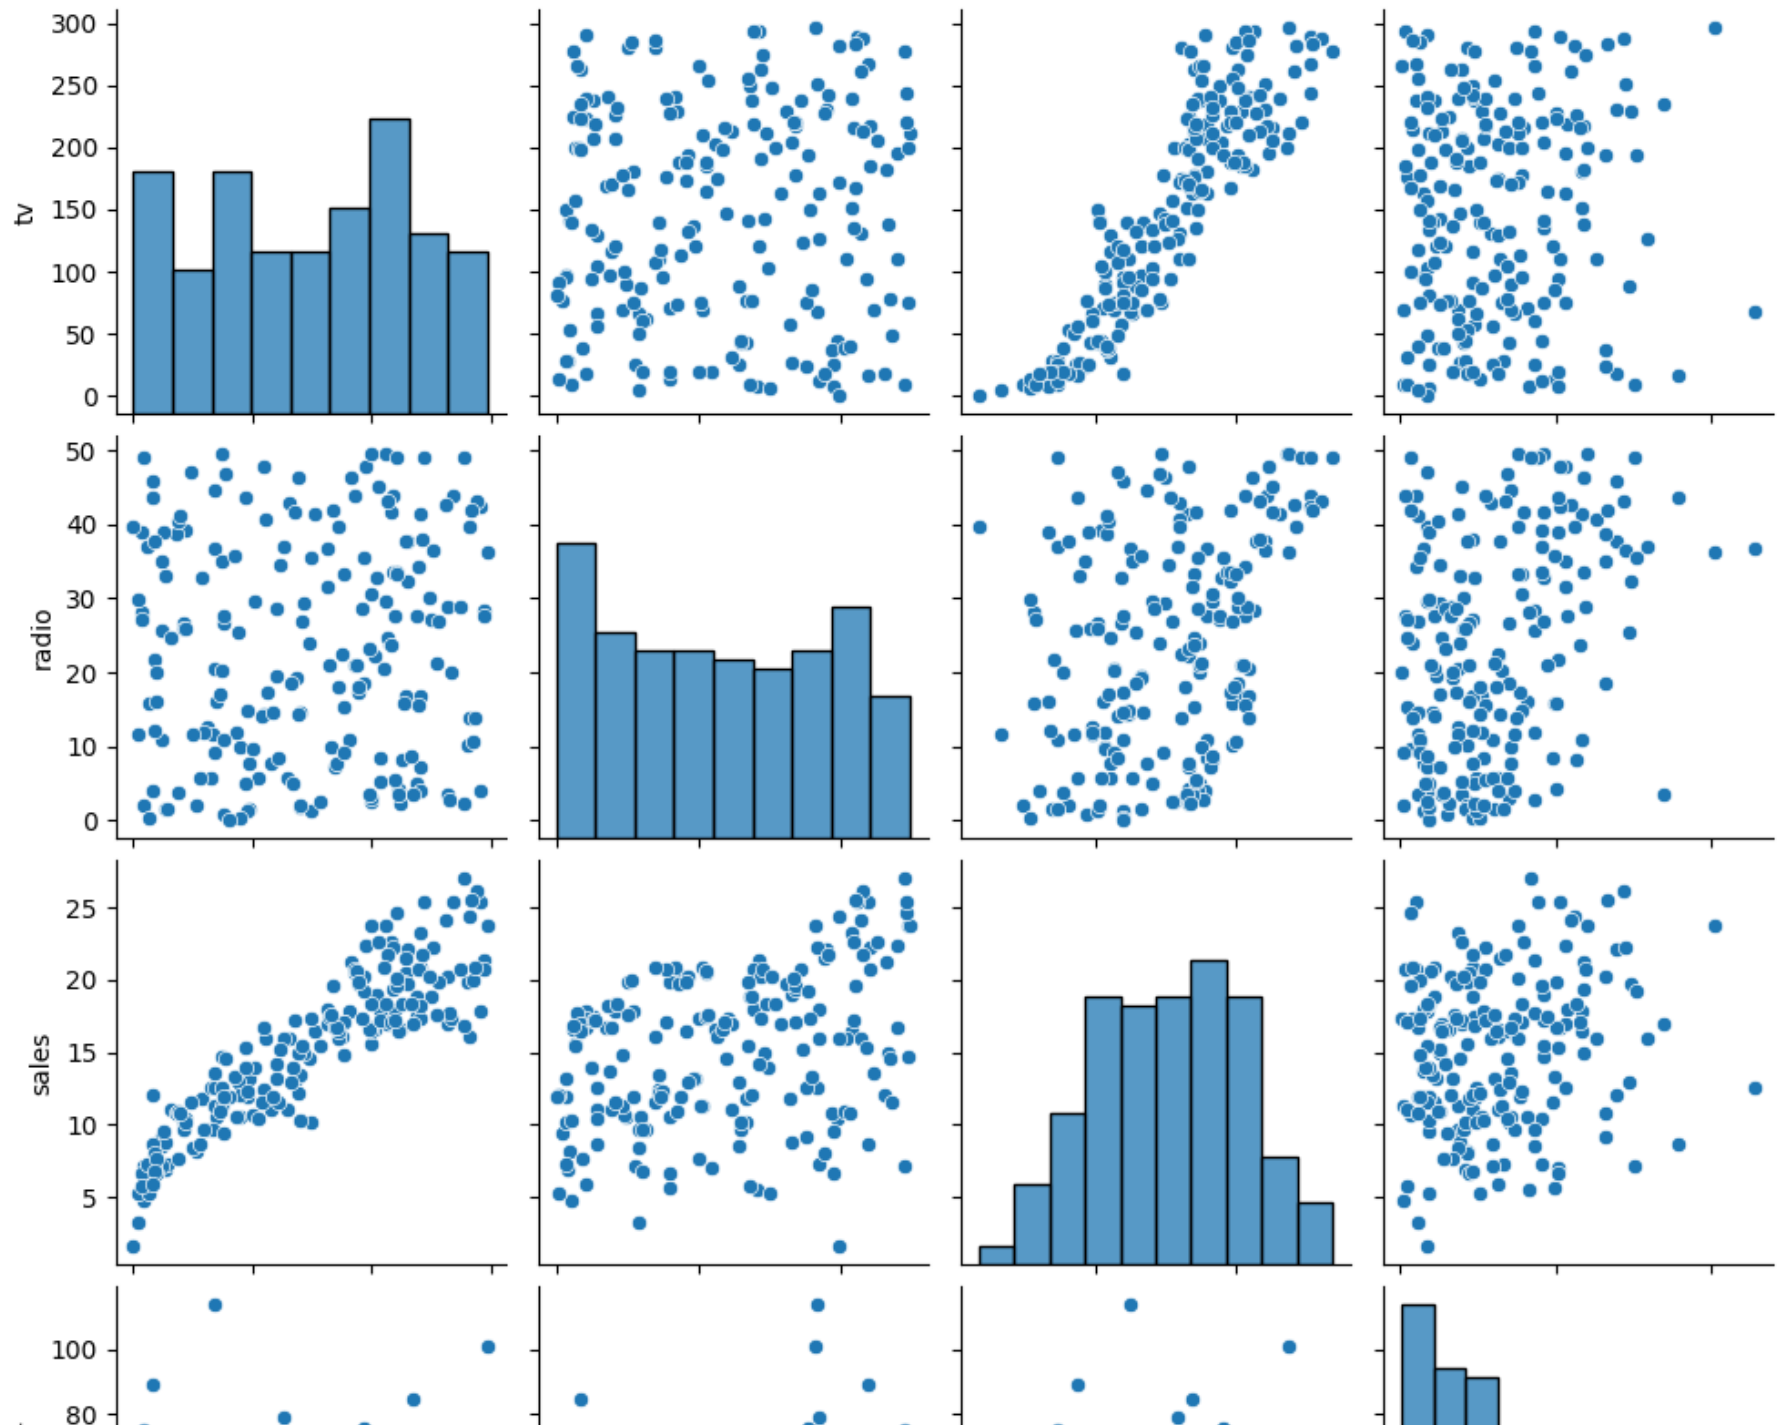
```
In [105]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.5)
```


```
In [106]: sns.pairplot(df)
```

```
Out[106]: <seaborn.axisgrid.PairGrid at 0x20e1f917580>
```











```
In [107]: sns.displot(df['sales'])
```

```
Out[107]: <seaborn.axisgrid.FacetGrid at 0x20e1f8c9990>
```




```
In [108]: sns.displot(df['tv'])
```

```
Out[108]: <seaborn.axisgrid.FacetGrid at 0x20e21a41a20>
```




```
In [109]: df500=df[:][50:500]  
sns.lmplot(x="radio",y="sales",data=df500,order=1)
```

```
Out[109]: <seaborn.axisgrid.FacetGrid at 0x20e1f9177f0>
```



```
In [110]: df500.fillna(method='ffill',inplace=True)
x=np.array(df['radio']).reshape(-1,1)
y=np.array(df['sales']).reshape(-1,1)
df.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(x_train,y_train)
print("Regression:",regr.score(x_test,y_test))
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='r')
plt.show()
```

Regression: 0.14284532571622577




```
In [111]: df.shape
```

```
Out[111]: (200, 4)
```

```
In [112]: df.isnull().sum()
```

```
Out[112]: tv          0  
radio          0  
sales          0  
newspaper      0  
dtype: int64
```


```
In [113]: sns.displot((y_test),bins=50);
```





```
In [114]: sns.pairplot(df,x_vars=['tv','radio','newspaper'],y_vars='sales',height=7,aspect=0.7,kind='reg')
```


```
Out[114]: <seaborn.axisgrid.PairGrid at 0x20e1e7ea1d0>
```



```
In [115]: plt.figure(figsize=(10,10))  
sns.heatmap(df.corr(),annot=True)
```

```
Out[115]: <Axes: >
```







```
In [116]: features = df.columns[0:2]
target = df.columns[-1]
#X and y values
x = df[features].values
y = df[target].values
#split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=17)
print("The dimension of x_train is {}".format(x_train.shape))
print("The dimension of x_test is {}".format(x_test.shape))
#Scale features
scaler = StandardScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.transform(x_test)
```

```
The dimension of x_train is (140, 2)
The dimension of x_test is (60, 2)
```

```
In [117]: lr = LinearRegression()  
#Fit model  
lr.fit(x_train, y_train)  
#predict  
#prediction = lr.predict(X_test)  
#actual  
actual = y_test  
train_score_lr = lr.score(x_train, y_train)  
test_score_lr = lr.score(x_test, y_test)  
print("\nLinear Regression Model:\n")  
print("The train score for lr model is {}".format(train_score_lr))  
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 0.08820570968809427

The test score for lr model is 0.17667256011867194

```
In [118]: ridgeReg = Ridge(alpha=10)  
ridgeReg.fit(x_train,y_train)  
#train and test scorefor ridge regression  
train_score_ridge = ridgeReg.score(x_train, y_train)  
test_score_ridge = ridgeReg.score(x_test, y_test)  
print("\nRidge Model:\n")  
print("The train score for ridge model is {}".format(train_score_ridge))  
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.08784976325183447


The test score for ridge model is 0.1684728204438981

```
In [120]: plt.figure(figsize=(10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge;$\alpha=10$')
plt.plot(Lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'Lasso;$\alpha=grid$')
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label=r'Linear Regression')
plt.xticks(rotation=90)
plt.grid()
plt.title("MY PROJECT")
plt.show()
```





## MY PROJECT





```
In [121]: print("\nLasso Model: \n")
lasso=Lasso(alpha=10)
lasso.fit(x_train,y_train)
train_score_ls=lasso.score(x_train,y_train)
test_score_ls=lasso.score(x_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```


Lasso Model:

The train score for ls model is 0.0

The test score for ls model is -0.0003547334659412815

```
In [122]: pd.Series(lasso.coef_,features).sort_values(ascending=True).plot(kind="bar")
```

```
Out[122]: <Axes: >
```

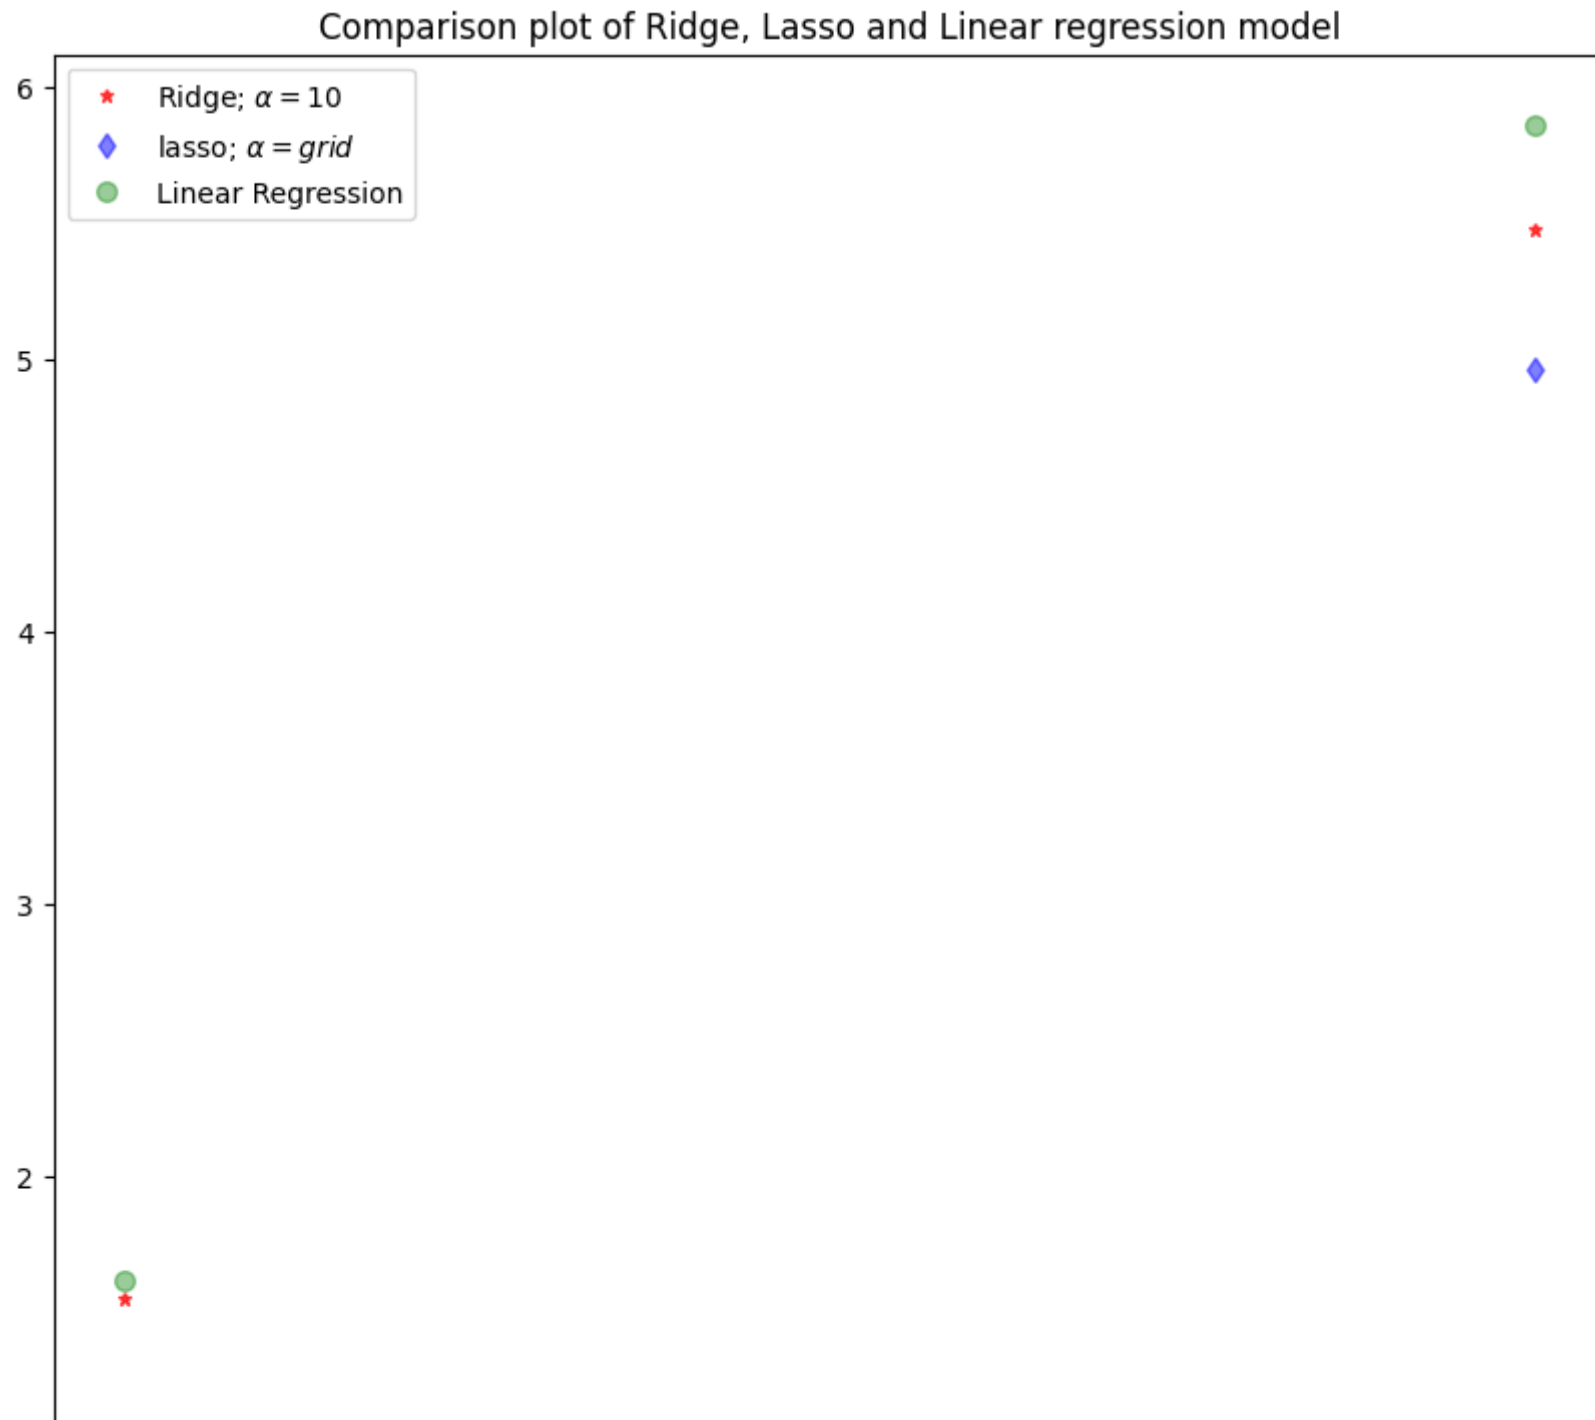


```
In [123]: from sklearn.linear_model import LassoCV
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,0.1,1,10],random_state=0).fit(x_train,y_train)
print(lasso_cv.score(x_train,y_train))
print(lasso_cv.score(x_test,y_test))
```

```
0.08414370280677297
0.16155807938746747
```

```
In [125]: #plot size
plt.figure(figsize = (10, 10))
#add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge; $\alpha = 0.7$')
#add plot for Lasso regression
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'lasso; $\alpha = 0.5$')
#add plot for linear model
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression; $\alpha = 0.4$')
#rotate axis
plt.xticks(rotation = 90)
plt.legend()
plt.title("Comparison plot of Ridge, Lasso and Linear regression model")
plt.show()
```







```
In [131]: from sklearn.linear_model import RidgeCV
ridge_cv=RidgeCV(alphas=[0.0001,0.001,0.01,0.1,1,10]).fit(x_train,y_train)
print("The train score for ridge model is {}".format(ridge_cv.score(x_train,y_train)))
print("The test score for ridge model is {}".format(ridge_cv.score(x_test,y_test)))
```

The train score for ridge model is 0.08784976325183436

The test score for ridge model is 0.16847282044389655

In [ ]: